


GENERATING MUSIC USING LSTM NEURAL NETWORKS


Author: Grzegorz Kazana

BIAI sem. 6


NEURAL NETWORKS AND SEQUENTIAL DATA




one to one




many to one




many to many




many to many



RNN

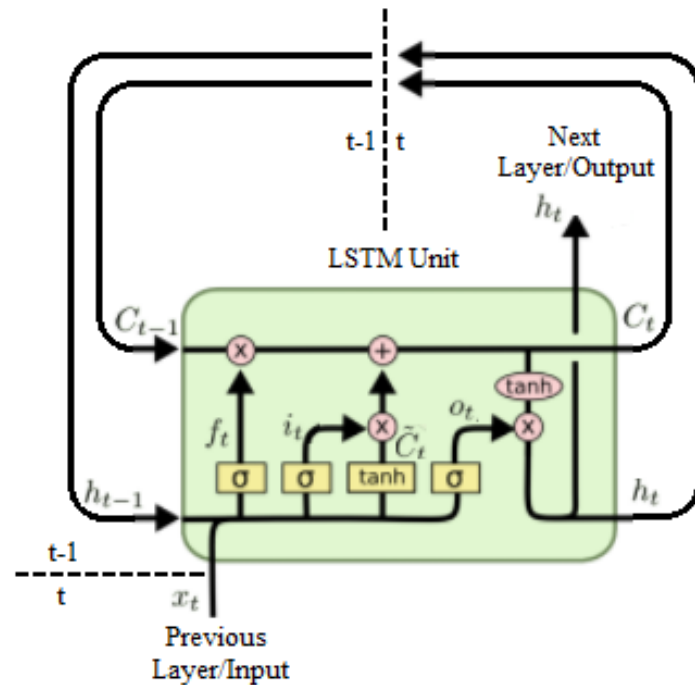


Recurrent Neural Network



LONG-SHORT TERM MEMORY NN

Understanding LSTM Networks



$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$


$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$


$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$


$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$






SEQUENTIAL DATA - MIDI




MIDI -> NUMPY

```
>>>
>>> for msg in mid.tracks[1][:100]:
...     print(msg)
...
<meta message track_name name='Piano right' time=0>
program_change channel=0 program=0 time=0
control_change channel=0 control=7 value=100 time=0
control_change channel=0 control=10 value=64 time=0
<meta message text text='bdca426d104a26ac9dcb070447587523' time=0>
control_change channel=0 control=91 value=127 time=0
note_on channel=0 note=60 velocity=35 time=2160
note_on channel=0 note=60 velocity=0 time=600
note_on channel=0 note=56 velocity=31 time=0
note_on channel=0 note=56 velocity=0 time=120
note_on channel=0 note=53 velocity=28 time=0
note_on channel=0 note=53 velocity=0 time=2160
note_on channel=0 note=56 velocity=35 time=0
note_on channel=0 note=56 velocity=0 time=600
note_on channel=0 note=60 velocity=30 time=0
note_on channel=0 note=60 velocity=0 time=120
note_on channel=0 note=65 velocity=35 time=0
note_on channel=0 note=65 velocity=0 time=720
note_on channel=0 note=68 velocity=35 time=0
note_on channel=0 note=68 velocity=0 time=600
note_on channel=0 note=72 velocity=30 time=0
```


```
>>>
>>> track = midi_to_numpy_pipe(mid)
>>> track
array([[0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       ...,
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.]])
>>>
>>> track.shape
(22355, 256)
>>>
>>> import matplotlib.pyplot as plt
>>> plt.imshow(track.T[:128, :600])
<matplotlib.image.AxesImage object at 0x00000165B30...
>>> plt.show()
```




A rectangular box with a black border, containing the text "X AND Y" in a bold, black, sans-serif font, centered within the box.




GENERATING NEW SAMPLES




LEARNING




Loss Function	Binary Crossentropy
Optimizer	Adam
Window size	50 (5s)
Hidden size	512
Samples per epoch	160000



Epoch #1	
Loss	0.0535
Accuracy	0.9823
False Positives	8513
False Negatives	109658




Epoch #2	
Loss	0.0380
Accuracy	0.9867
False Positives	27069
False Negatives	65661




Epoch #3	
Loss	0.0380
Accuracy	0.9867
False Positives	27069
False Negatives	65661



Epoch #4	
Loss	0.0284
Accuracy	0.9893
False Positives	23349
False Negatives	52003





Epoch #5	
Loss	0.0231
Accuracy	0.9914
False Positives	19882
False Negatives	43153





Epoch #6	
Loss	0.0194
Accuracy	0.9928
False Positives	14648
False Negatives	33459




Epoch #7	
Loss	0.0145
Accuracy	0.9951
False Positives	9852
False Negatives	23792






Epoch #10	
Loss	0.0065
Accuracy	0.9982
False Positives	3998
False Negatives	7885



Epoch #14	
Loss	0.0036
Accuracy	0.9989
False Positives	3079
False Negatives	4675



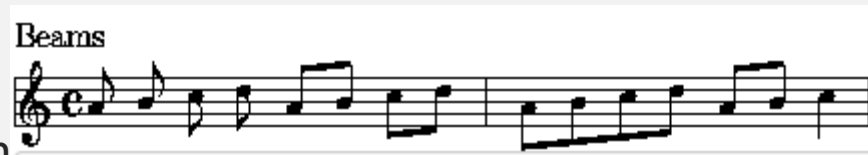
Epoch #20	
Loss	0.0027
Accuracy	0.9991
False Positives	2404
False Negatives	3526



TO-DO

- Use more data
- Rethink Midi en
- Tune hyper-para
- Test different w
- Write custom loss function penalizing false negatives
- Add more layers
- Use music sheet .abc format

Beams



```
X:1  
T:Beams  
M:C  
K:C  
A B c d AB cd|ABcd ABc2|]
```

SOURCES

- <https://towardsdatascience.com/back-to-basics-deriving-back-propagation-on-simple-rnn-lstm-feat-aidan-gomez-c7f286ba973d>
- <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>
- <https://medium.com/datadriveninvestor/how-do-lstm-networks-solve-the-problem-of-vanishing-gradients-a6784971a577>

A screenshot of a Jupyter Notebook interface. The top bar shows the title "Deep Learning for Computer Vision.ipynb" and various icons for file operations, comments, and sharing. Below the title bar is a menu with "File", "Edit", "View", "Insert", "Runtime", "Tools", and "Help". The main area of the notebook is divided into two sections: a code cell and a text cell. The code cell contains the following Python code:

```
[ ] :om keras import layers
      :om keras import models

      >del = models.Sequential()
      >del.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28
      >del.add(layers.MaxPooling2D((2, 2)))
      >del.add(layers.Conv2D(64, (3, 3), activation='relu'))
      >del.add(layers.MaxPooling2D((2, 2)))
      >del.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

The text cell below the code cell contains the following text:

A convnet is a stack of Conv2D and MaxPooling2D layers. A convnet takes as input tensors of shape (image_hight, image_width, image_channels).